

Lecture 26 — Finding Bottleneck Devices

Jeff Zarnett

jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

September 16, 2024



Legos came broken

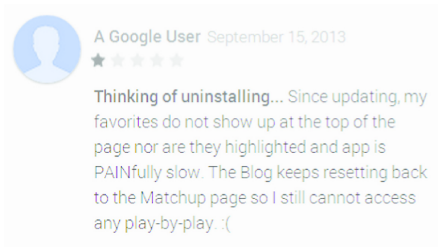
Reviewed on December 23, 2019

Style: Standard | **Verified Purchase**

My Lego set was broken when I opened it!



Poor mobile app performance is a major complaint in app store reviews.



We might have a *vague* idea what's wrong, but how do we fix it?

Who Can It Be Now?



We usually assume that CPU is the problem... but is that true?

Future topics will mostly follow CPU profiling, but it's not the only thing.

Elementary, My Dear Watson!

It is a capital mistake to theorize before one has data. Insensibly one begins to twist facts to suit theories, instead of theories to suit facts.

— Sherlock Holmes
(*A Scandal in Bohemia*; Sir Arthur Conan Doyle)



Who's to blame?

- 1 CPU
- 2 Memory
- 3 Disk
- 4 Network
- 5 Locks

These are, obviously,
categories rather than specific causes.

Fixing it might involve using techniques from this course...
But it might be really difficult.

Is the user's hardware just too old?

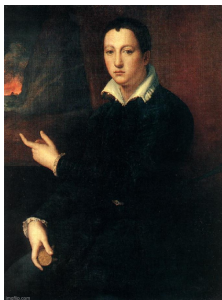


The “slow” workflow might actually be a bug...

Are you doing a slow or async task in the UI thread?

Oops! Just fix the bug.

CPU is probably the easiest of these to diagnose.



htop, Task Manager, etc. will tell you if CPU hosed.

Look at the %CPU columns and see where all your CPU is going.

Still, that tells you about right now; what about the long term average?

Checking with my machine “Loki”, which has since ascended to Valhalla:

```
top - 07:28:19 up 151 days, 23:38, 8 users, load average: 0.87, 0.92, 0.91
```

Those last three numbers are the 1, 5, and 15 minute averages of CPU load.

Lower numbers mean less CPU usage and a less busy machine.

Picture a single core of a CPU as a lane of traffic.

You are a bridge operator and so you need to monitor how many cars are waiting to cross that bridge.

If no cars are waiting, traffic is good and drivers are happy.

If there is a backup of cars, then there will be delays.

- 1** 0.00 means no traffic.
Anything between 0.00 and 0.99 means we're under capacity and there will be no delay.
- 2** 1.00 means we are exactly at capacity.
Everything is okay, but if one more car shows up, there will be a delay.
- 3** Anything above 1.00 means there's a backup (delay).
If we have 2.00 load, then the bridge is full and there's an equal number of cars waiting to get on the bridge.



= load of 1.00



= load of 0.50



= load of 1.70

≥ 1.00 isn't necessarily bad, but you should be concerned if there is consistent load of 1.00 or above.

< 1.00 but getting close to it: you know how much room you have to scale things up.

> 0.70 then it's probably time to investigate.

≥ 1.00 consistently we have a serious problem.

5.00: this is a red alert situation.

Now this is for a single CPU—if you have a load of 3.00 and a quad core CPU, this is okay.

Traffic analogy: four lanes of traffic, of which 3 are being used to capacity.

So we have a 4th lane free and it's as if we're at 75% utilization on a single CPU.

Using garbage collection? Is it running a lot?



Out of memory errors (crash or recovery)?

How to tell? Look at disk utilization.

Not enough RAM \Rightarrow swapping, bad perf, no scalability.

(In the worst case.)

You can ask via `top` about how much swap is being used, but that's probably not the interesting value.

```
KiB Mem:   8167736 total,  6754408 used,  1413328 free,   172256 buffers
KiB Swap:  8378364 total,  1313972 used,  7064392 free. 2084336 cached Mem
```

Why? Memory being “full” does not necessarily mean anything bad.

It means the resource is being used to its maximum potential, yes, but there is no benefit to keeping a block of memory open for no reason. (Or stockpiling late days).

Also, memory is unlike CPU; if there's nothing for the CPU to do, it will just idle (low power state).

Memory won't "forget" data if it doesn't happen to be needed right now—data will hang around in memory until there is a reason to move or change it.

So freaking out about memory appearing as full is kind of like getting all in a knot about how "System Idle Process" is hammering the CPU...

You can also ask about page faults, with the command

```
ps -eo min_flt,maj_flt,cmd
```

Major page faults: had to fetch from disk.

Minor page faults: had to copy a page from another process.

The output of this is too big even for the notes.

This is process-lifetime data.

What you really want is to ask Linux for a report on swapping:

```
jz@Loki:~$ vmstat 5
```

procs		-----memory-----				---swap--		-----io-----		-system--			-----cpu-----			
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st
1	0	1313972	1414600	172232	2084296	0	0	3	39	1	1	27	1	72	0	0
0	0	1313972	1414476	172232	2084296	0	0	0	21	359	735	19	0	80	0	0
0	0	1313972	1414656	172236	2084228	0	0	0	102	388	758	22	0	78	0	0
4	0	1313972	1414592	172240	2084292	0	0	0	16	501	847	33	0	67	0	0
0	0	1313972	1412028	172240	2084296	0	0	0	0	459	814	29	0	71	0	0

Swapping Report with Actual Swapping

procs			swpd	free	memory			swap			io		system		cpu
r	b	w			buff	cache	si	so	bi	bo	in	cs	us	sy	
.	.	.													
1	0	0	13344	1444	1308	19692	0	168	129	42	1505	713	20	11	
1	0	0	13856	1640	1308	18524	64	516	379	129	4341	646	24	34	
3	0	0	13856	1084	1308	18316	56	64	14	0	320	1022	84	9	

Looking at disk might seem slightly redundant if memory is not the limiting factor.

After all, if the data were in memory it would be unnecessary to go to disk in the first place.

Still, sometimes we can take a look at the disk and see if that is our bottleneck.

```
jz@Loki:~$ iostat -dx /dev/sda 5
Linux 3.13.0-24-generic (Loki) 16-02-13 _x86_64_ (4 CPU)
```

Device:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	r_await	w_await	svctm	%util
sda	0.24	2.78	0.45	2.40	11.60	154.98	116.91	0.17	61.07	11.57	70.27	4.70	1.34

Last column %util tells us what we want to know.

We can ask about the network with `nload`.

You get a nice little graph if there is anything to see.

But you'll get the summary, at least:

```
Curr: 3.32 kBit/s  
Avg: 2.95 kBit/s  
Min: 1.02 kBit/s  
Max: 12.60 kBit/s  
Ttl: 39.76 GByte
```

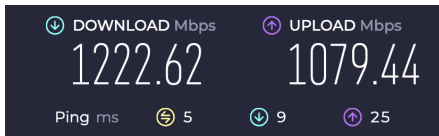
If we saw data leaving at 100 MBit, might be clear it's the max.

But what if you're using other hardware?



Same for wireless networks... walls, floors, EM interference...

How to test? Use tools like speedtest.net



Need to test multiple times.

Good results doesn't necessarily mean good performance.

If you want to get an idea of the path and the latency to a particular remote system, you can use the `tracert` tool.

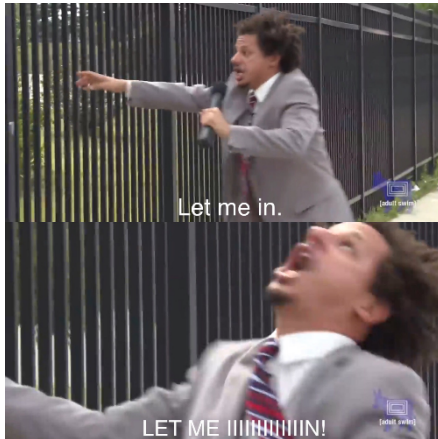
```
Microsoft Windows [Version 10.0.19043.1288]
(c) Microsoft Corporation. All rights reserved.
C:\Users\Michael>tracert catchpoint.com
Tracing route to catchpoint.com [64.79.149.76]
Over a maximum of 30 hops:
 1  2ms  1ms  1ms  10.0.0.1
 2  10ms 10ms 10ms 96.120.40.245
 3  10ms 11ms 12ms 96.110.175.85
 4  10ms 16ms 10ms 162.151.63.57
 5  19ms 16ms 20ms 96.108.21.57
 6  15ms 19ms 14ms 96.216.134.10
 7      19ms 22ms 21ms be-32121-cs02.350ecermak.il.ibone.comcast.net [96.110.42.181]
 8      22ms 34ms 22ms be-2204-pe04.350ecermak.il.ibone.comcast.net [96.110.37.38]
 9      22ms 20ms 20ms 50.208.234.106
10     51ms 50ms 49ms ae18-0.cr02.dl1s02-tx.us.windstream.net [40.128.10.135]
11     73ms 72ms 72ms ae4-0.agr03.phnd01-az.us.windstream.net [169.130.193.231]
12     84ms 73ms 75ms ae1-0.pe05.phnd01-az.us.windstream.net [169.130.169.31]
13     85ms 84ms 85ms h241.23.132.40.static.ip.windstream.net [40.132.23.241]
14      *      82ms 78ms be181.las-n10s1-core1.switch.com [66.209.64.121]
15     79ms 77ms 80ms bell011.las-agg7s5-1.switch.com [66.209.72.26]
16     79ms 77ms 79ms 64.79.139.18
17     77ms 77ms 87ms 64.19.149.76
Trace complete
```

Latency can never get down to 0: speed of light limitation.

Example: New York to Lyon is 73.21ms which is something like 83.79% of the speed of light in fibre-optic cable (as of August 2024).

Another cause: packet loss! Requires re-sending data...
Replace dying devices? Environmental causes?

Maybe our code is slow because we are waiting for locks?



Out of scope: deadlock detection.

Unexpectedly low CPU usage not explained by I/O-waiting?

Many threads blocked?

No magical lock trace tool – may need our own tracing.

perf lock is for kernel locks; VTune costs money.

Most of our discussion will be about CPU though.

When you finally discover who's been sabotaging you all along



Why? That's probably what it is!